

# Peer-to-peer protocols for efficient and resilient decentralized learning


Mohamed Amine Legheraba

Supervisors: Maria Potop-Butucaru, Sébastien Tixeul

LIP6 · Sorbonne Université

September 7, 2026


# Context



# Supervised Classification

## Training

Data




Dog



Cat



(data, label)



Model

Labels


## Prediction



Trained Model

Cat

# Federated Learning



---

## Server

---

```
1 loop  
2   received_models  $\leftarrow$  collect(clients)  
3   global_model  $\leftarrow$  merge(received_models)  
4   broadcast(global_model)
```

---

---

## Client



---

```
1 loop  
2   local_model  $\leftarrow$  train(model, local_data)  
3   send(local_model, server)  
4   model  $\leftarrow$  receive(global_model)
```


---

[1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Artificial intelligence and statistics*, 2017, pp. 1273–1282.

# Structured vs Unstructured networks



# Gossip Learning



---

## Gossip Learning (main thread)

---

ML model: **model**


List of neighbors : **cache**

```
1 loop
2   | Wait for  $\Delta$  time units
3   |  $\text{peer} \leftarrow \text{selectRandom}(\text{cache})$ 
4   |  $\text{send}(\text{peer}, \text{model})$ 
```

---

[2] R. Ormándi, I. Hegedűs, and M. Jelasity, “Gossip learning with linear models on fully distributed data,” *Concurrency and Computation: Practice and Experience*, vol. 25, no. 4, pp. 556–571, 2013.

# Gossip Learning vs Federated Learning




[3] I. Hegedűs, G. Danner, and M. Jelasity, “Decentralized learning works: An empirical comparison of gossip learning and federated learning,” *Journal of Parallel and Distributed Computing*, vol. 148, pp. 109–124, 2021.

# State of the Art Summary

|                                             | Decentralized | Local data | Fast convergence           | Fault-tolerant | No overlay overhead |
|---------------------------------------------|---------------|------------|----------------------------|----------------|---------------------|
| Centralized learning                        | ✗             | ✗          | ✓                          | ✗              | ✓                   |
| Federated learning                          | ✗             | ✓          | ✓                          | ✗              | ✓                   |
| Decentralized learning<br>(structured net.) | ✓             | ✓          | ✓                          | ✗              | ✗                   |
| Gossip learning                             | ✓             | ✓          | ✗                          | ✓              | ✓                   |
| Local learning                              | ✓             | ✓          | ✗<br>no generalizing model | ✓              | ✓                   |

**No single solution is fully satisfactory**





## Outline

- **Elevator** — overlay (hub election)
- **Lift** — malicious resilience
- **HEAL** — model aggregation
- **FLAIR** — wireless networks adaptation

# Elevator

---


**Definition.** (Hub) Let  $G = (V, E)$  be the directed overlay graph, where each node  $v \in V$  keeps a partial view  $P(v) \subseteq V$ .

A node  $h \in V$  is a **hub** if it appears in the partial view of every node:


$$\forall v \in V, \quad h \in P(v)$$

[4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.


1. One node is more connected





2. It attracts more connections



3. Snowball effect: it becomes a hub



Preferential attachment — Barabási & Albert [5] A.-L. Barabási, H. Jeong, Z. Néda, E. Ravasz, A. Schubert, and T. Vicsek, “Evolution of the social network of scientific collaborations,” *Physica A: Statistical mechanics and its applications*, vol. 311, no. 3–4, pp. 590–614, 2002.



Before and after a shuffling operation. Node 1 sends addresses {itself, 2, 3} to node 4. Node 4 sends back {5,6, 8}.

[6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.

**Preferential attachment** makes hubs emerge...

**Random attachment** keeps the network resilient...

Can we combine both to get **hubs** that are  
**resilient**?

## Mechanism

1. Each node asks its **neighbours** for **their** neighbour lists.
2. It connects to the  **$h$  most frequent** nodes.
3. It asks these  $h$  hubs for extra incoming connections, to fill a **random subset** of its view.

## Observed

- After a few cycles, **hubs emerge** (the preferred  $h$  are the same for everyone);
- Hubs are **elected at random** among the network nodes;
- Even after **hub failures**, new nodes are elected as hubs.



# Example of topology generated





Figure 1: Indegree distribution of a network generated with Elevator, with 1000 nodes and 10 hubs



**Definition.** (Stability) Once Elevator has converged to a set of  $h$  hubs, both the list of hubs and their number  $h$  remain constant (with high probability) over time.


Formally, after convergence time  $T$ :

$$\Pr[\forall t \geq T, H(t) = H(T) \wedge \forall v \in V, C_v(t)[1 : h] = H(t)] = 1$$

**Definition.** (Convergence) The network converges to a stable state containing **exactly**  $h$  hubs shared by all nodes.

## Idea of the proof.

- A new hub is produced with **non-zero probability** at each cycle.
- The number of hubs **cannot decrease** (until it reaches  $h$ ).
- Once at least one hub exists, the network is **strongly connected** (w.h.p.) and produces a new hub.



PeerSim was **forked and substantially rewritten** for this thesis:

- Migrated from **SVN to Git**, build rewritten in **Gradle**;
- **Docker + GitLab CI/CD**;
- **Parallelised** the cycle-based engine;
- Failure models from scratch (**crash, churn, Byzantine**);

## Parameters

- Network size:  **$N = 1000$  nodes**;
- **1000 cycles**, repeated **100 times**;
- Initial topology: **random k-out graph** with  **$k = c = 20$** ;
- **$h = 10$  hubs**;
- e.g. brutal crash: **50% of nodes** disconnected at cycle **500**.

## Overlay

- In-degree distribution
- Clustering coefficient
- Average path length
- Diameter
- Convergence time

## Machine learning



- Accuracy

## Overlay



- **Crash failures**
- **Churn** (dynamic joins / leaves)
- **Byzantine failures**

## Machine learning




- **Privacy attacks**
- **Poisoning attacks**

**Note:** ML-level failures are not addressed in this thesis



ONNE  
RSITE

- **Go + libp2p** over **TCP/IP streams** (+ light **HTTP** per node);
- networks of **20 to 50 nodes** on **1–2 machines**;
- modes: **synchronous / externally-synchronized / asynchronous**;
- **hubs emerge within the first few cycles**  
— matching theory & simulation.





# Lift

---

---

Coordinated attack — lying about neighbors

---


addresses of all colluding nodes: **colluders**

backward list: **backward\_peers**

```
1 loop
2   (request, peer) ← receive()
3   if request == CACHE_REQUEST
4     backward_peers.add(peer)
5     colluders.shuffle()
6     modified_cache ← colluders[0:c]
7     send(modified_cache, peer)
```

---

Each colluding node answers with a fake cache that only contains **other colluding nodes**, inflating their visibility.



After convergence, correct nodes **deterministically** re-draw the hubs from a shared random seed — attackers cannot bias the outcome.

---

## Hub redistribution


---

```
current hub list: H
network size: N
target hubs: h
1 seed  $\leftarrow$  sort(H)
2 prng  $\leftarrow$  Random(seed)
3 selected  $\leftarrow$  {}
4 while selected.size() < h
5   | id  $\leftarrow$  prng.nextInt(N)
6   | if id not in selected
7   |   | selected  $\leftarrow$  selected  $\cup$  {id}
8 H  $\leftarrow$  selected
```

---


# Resilience against colluding attackers

Evolution of number of hubs (100 cycles)




HEAL





1) Each node train a local model

2) Each node sends its model to a (random) hub



3) Hubs receive and aggregate models

4) Hubs compute a global model together

5) Hubs send back the global model to the nodes

[7] M. A.

Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeul, “HEAL: Resilient and Self-\* Hub-based Learning,” in *International Conference on Advanced Information Networking and Applications*, 2025, pp. 200–211.




## Spambase + Logistic regression

Variables Table

| Variable Name     | Role    | Type       |
|-------------------|---------|------------|
| word_freq_make    | Feature | Continuous |
| word_freq_address | Feature | Continuous |
| word_freq_all     | Feature | Continuous |
| word_freq_3d      | Feature | Continuous |
| word_freq_our     | Feature | Continuous |
| word_freq_over    | Feature | Continuous |
| word_freq_remove  | Feature | Continuous |

4601 emails, binary  
57 parameters


## MNIST + LeNet5





70k digits, 10 classes  
CNN, 60k parameters

# Context of experiments

- **Peer-to-Peer:** On the PeerSim simulator.
  - Each generated network had 100 nodes, and the Hub-based topology had 5 hubs
  - In addition to the crash-free context, we ran the algorithm during a context of crash, a context of churn and a context of an attack on the hubs
- **Decentralized Learning:** On the gossip simulator
  - **ML models:** Logistic Regression & LeNet5
  - **Datasets:** Spambase and MNIST
  - Comparison with Federated Learning, Gossip Learning, Epidemic Learning, GAIA, Chord-based Learning, Fedlay

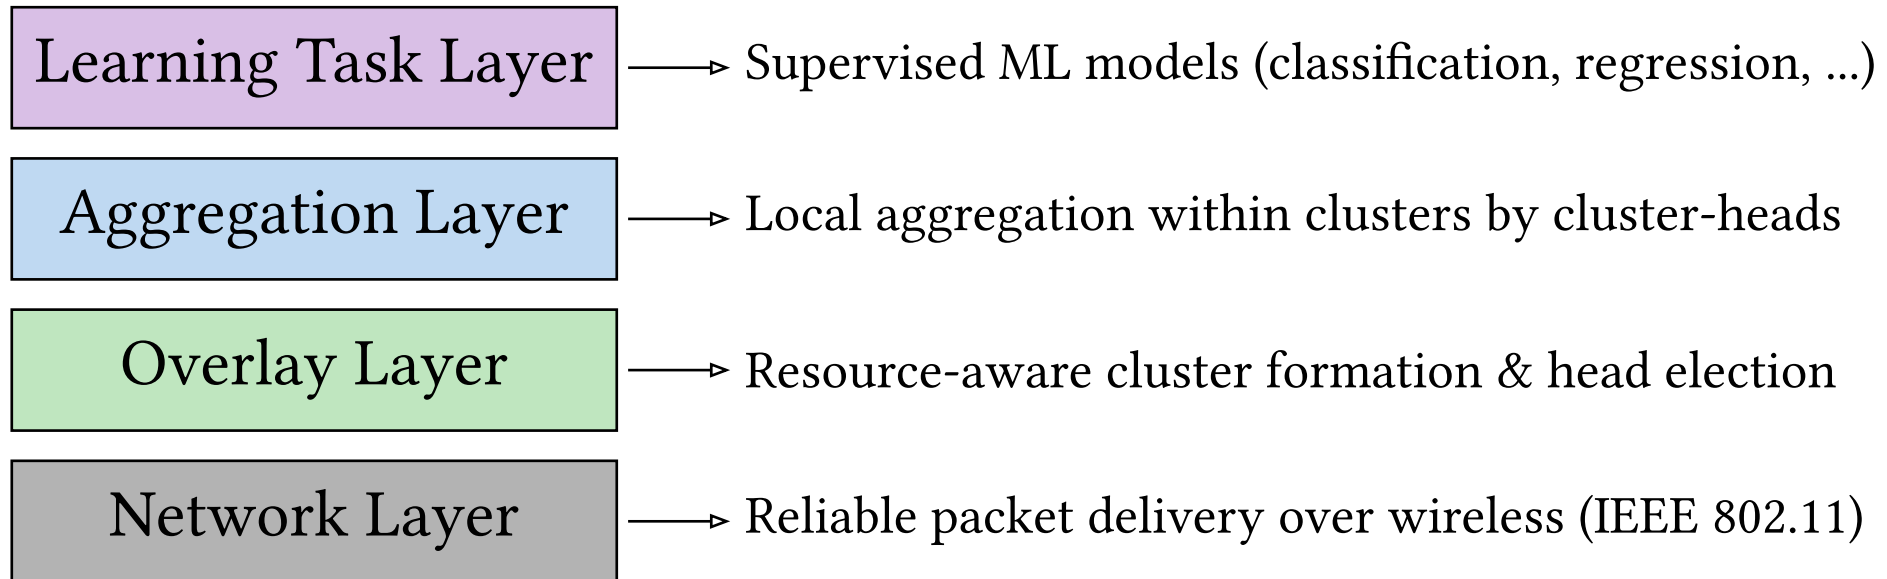






# FLAIR

---



FLAIR operates in **rounds** (clustering inspired by LEACH); a fraction  $p \in (0, 1)$  of nodes serve as **cluster-heads** (CHs) each round, and the role **rotates** to balance the load.

A node eligible (not a CH during the last  $1/p$  rounds) elects itself if  $x < T(n)$ , with  $x \sim \text{Uniform}(0, 1)$  drawn via a **VRF**:


$$T(n) = p \cdot R_n / (1 - p \cdot (r \bmod 1/p))$$

with  $R_n = \alpha \cdot \text{CPU}_n + \beta \cdot \text{RAM}_n + \gamma \cdot \text{GPU}_n + \delta \cdot \text{BW}_n$  (resource score;  $\alpha + \beta + \gamma + \delta = 1$ ).

Nodes join the nearest CH; the CH **aggregates locally** within its cluster.



- Static network (100 nodes): **highest accuracy**  $\approx 0.91$  (C-FL / HEAL  $\approx 0.90$ , Gossip  $\approx 0.88$ )
- Convergence up to **2.5x faster** than Gaia
- Resilient to **permanent / temporary / random crashes** (up to 90% nodes)
- Tested under **mobility** and in a **smart farming** scenario



# Conclusion

---

- **Elevator**: a self-organising overlay that makes hubs emerge — provides a **structured overlay** for structured aggregation without any coordinator.
- **HEAL**: federated learning directly onto the structured overlay — recovers the efficiency of FL while staying decentralised.
- **Lift**: hardens hub election against colluding byzantines (up to 10%).
- **FLAIR**: validates the modular architecture in wireless networks (LEACH-style clustering).

**Answer**: structured aggregation is possible **without centralisation**, if the overlay layer is designed with that goal in mind.

- **Capability-aware hub election** (heterogeneous nodes) and sensitivity to the initial topology.
- Optimise Elevator (message complexity, memory, simplicity).
- Address **data heterogeneity** in HEAL (personalisation, locally-weighted aggregation).
- A unified **Python simulator** for peer sampling + ML.

This work has already begun during a **3-month internship at NII (Tokyo)**.

- Better robustness at the Overlay layer (using)
- **Byzantine robustness in the learning layer** (poisoning & privacy attacks)
- Extend HEAL to **unsupervised and reinforcement learning**.
- **Vertical federated learning** (different features per participant).
- **Formal convergence guarantees** for HEAL.

- [8] M. A. Legheraba, M. Potop-Butucaru, and S. Tixeul, “Brief Announcement: A Self-\* and Persistent Hub Sampling Service,” in *International Symposium on Stabilizing, Safety, and Security of Distributed Systems*, 2024, pp. 461–465.
- [4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeul, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.
- [9] M. A. Legheraba, N. Rachdi, M. Potop-Butucaru, and S. Tixeul, “LIFT: Byzantine Resilient Hub-Sampling,” in *2025 Thirteenth International Symposium on Computing and Networking (CANDAR)*, 2025, pp. 1–9. — **Outstanding Paper Award**
- [7] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeul, “HEAL: Resilient and Self-\* Hub-based Learning,” in *International Conference on Advanced Information Networking and Applications*, 2025, pp. 200–211.
- [10] I. Boutebicha, B. Zaghdoudi, M. A. Legheraba, and M. P. Butucaru, “FLAIR: Distributed Federated Learning with Dynamic Clustering,” in *Proceedings of the 14th International Conference on Networked Systems (NETYS 2026)*, in Lecture Notes in Computer Science. Agadir, Morocco: Springer, 2026.
- [11] M. A. Legheraba, M. Potop-Butucaru, S. Tixeul, and S. Fdida, “Des noeuds anonymes hissés au rang de stars,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- [12] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeul, “Les étoiles montantes de l'apprentissage distribué,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- A journal extension (journal version of this work) has been **submitted to the IEEE/ACM Transactions on Networking**; we are awaiting the **final review** (currently in **minor revision**).

## Bibliography

- [1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Artificial intelligence and statistics*, 2017, pp. 1273–1282.
- [2] R. Ormándi, I. Hegedűs, and M. Jelasity, “Gossip learning with linear models on fully distributed data,” *Concurrency and Computation: Practice and Experience*, vol. 25, no. 4, pp. 556–571, 2013.
- [3] I. Hegedűs, G. Danner, and M. Jelasity, “Decentralized learning works: An empirical comparison of gossip learning and federated learning,” *Journal of Parallel and Distributed Computing*, vol. 148, pp. 109–124, 2021.
- [4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.
- [5] A.-L. Barabási, H. Jeong, Z. Néda, E. Ravasz, A. Schubert, and T. Vicsek, “Evolution of the social network of scientific collaborations,” *Physica A: Statistical mechanics and its applications*, vol. 311, no. 3–4, pp. 590–614, 2002.
- [6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.
- [7] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeuil, “HEAL: Resilient and Self-\* Hub-based Learning,” in *International Conference on Advanced Information Networking and Applications*, 2025, pp. 200–211.

- [8] M. A. Legheraba, M. Potop-Butucaru, and S. Tixeuil, “Brief Announcement: A Self-\* and Persistent Hub Sampling Service,” in *International Symposium on Stabilizing, Safety, and Security of Distributed Systems*, 2024, pp. 461–465.
- [9] M. A. Legheraba, N. Rachdi, M. Potop-Butucaru, and S. Tixeuil, “LIFT: Byzantine Resilient Hub-Sampling,” in *2025 Thirteenth International Symposium on Computing and Networking (CANDAR)*, 2025, pp. 1–9.
- [10] I. Boutebicha, B. Zaghdoudi, M. A. Legheraba, and M. P. Butucaru, “FLAIR: Distributed Federated Learning with Dynamic Clustering,” in *Proceedings of the 14th International Conference on Networked Systems (NETYS 2026)*, in Lecture Notes in Computer Science. Agadir, Morocco: Springer, 2026.
- [11] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Des noeuds anonymes hissés au rang de stars,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- [12] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeuil, “Les étoiles montantes de l'apprentissage distribué,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- [13] A. Montresor and M. Jelasity, “PeerSim: A Scalable P2P Simulator,” in *Proc. of the 9th Int. Conference on Peer-to-Peer (P2P'09)*, Seattle, WA, Sep. 2009, pp. 99–100.
- [14] M. Jelasity, S. Voulgaris, R. Guerraoui, A.-M. Kermarrec, and M. Van Steen, “Gossip-based peer sampling,” *ACM Transactions on Computer Systems (TOCS)*, vol. 25, no. 3, p. 8-es, 2007.




- [15] J. Huang, L. Kong, G. Chen, Q. Xiang, X. Chen, and X. Liu, “Blockchain-Based Federated Learning: A Systematic Survey,” *IEEE Network*, vol. 37, no. 6, pp. 150–157, 2023.



# Appendix



---



## Federated Learning (FedAvg)

```
1 for  $t = 0$  to  $T - 1$  do
2   server broadcasts  $\theta^{(t)}$ 
   for each client  $i$  in parallel:
3      $\theta_i \leftarrow$  local training (lr  $\eta$ ,  $E$ 
       steps)
4     server aggregates:  $\theta^{(t+1)} \leftarrow$ 
        $\sum w_i \theta_i$ 
5 end for
```

[1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Artificial intelligence and statistics*, 2017, pp. 1273–1282.



---

## HEAL Learning: The Hub Algorithm

---

```
duration to wait for model: delta_time
list of all hubs: hubs_list
1 nb_hubs  $\leftarrow$  hubs_list.size()
2 loop
3   models  $\leftarrow$  {}
4   time  $\leftarrow$  time.now()
5   backwards_list  $\leftarrow$  {}
6   while time.now() < time + delta_time do
7     (peer_model, peer)  $\leftarrow$  receive()
8     models.append(peer_model)
9     backwards_list.append(peer)
10  average_model  $\leftarrow$  average(models)
11  send(hubs_list, average_model)
12  models_hubs  $\leftarrow$  {}
13  models_hubs.append(average_model)
14  nb_receive  $\leftarrow$  0
15  while nb_receive < nb_hubs - 1 do
16    hub_model  $\leftarrow$  receive()
17    nb_receive  $\leftarrow$  nb_receive + 1
18    models_hubs.append(hub_model)
19  global_model  $\leftarrow$  average(models_hubs)
20  send(backwards_list, global_model)
```

---

---

## HEAL Learning: The Client Algorithm

---

```
duration to wait for model: delta_time
The ML model, initialized at random: model
The local data: data
list of all hubs: hubs_list
number of hubs to send the model: s
1 loop
2   model  $\leftarrow$  trainModel(model, data)
3   hubs  $\leftarrow$  chooseRandom(hubs_list, s)
4   for hub in hubs do
5     | send(hub, model)
6   hubs_models  $\leftarrow$  {}
7   for hub in hubs do
8     | model_hub  $\leftarrow$  receive()
9     | hubs_models.append(model_hub)
10  model  $\leftarrow$  average(hubs_models)
```

---

---

### Elevator Algorithm (background thread)

---

```
1 loop
2   request, peer  $\leftarrow$  receive()
3   if request == CACHE_REQUEST then
4     send(cache, peer)
5     backward_peers.add(peer)
6   if request == BACKWARD_REQUEST then
7     send(backward_peers, peer)
```

---

- Simulations done with the Java PeerSim<sup>\*</sup> simulator (modified),  
with the cycle based mode
- Comparaisons against 3 peer sampling algorithms : Newscast, Proofs and Phenix (Power-law)



|                                                               | Value |
|---------------------------------------------------------------|-------|
| Network size                                                  | 1000  |
| Number of cycles for each simulation                          | 1000  |
| Number of times each simulation was run (with different seed) | 100   |
| c parameter (cache size)                                      | 20    |
| h parameter (number of hubs)                                  | 10    |

- All simulations were run on 16 vCPU, using 64G of memory.
- Code is available at <https://gitlab.lip6.fr/legheraba/elevator>



---



\*



[13] A. Montresor and M. Jelasity, "PeerSim: A Scalable P2P Simulator," in *Proc. of the 9th Int. Conference on Peer-to-Peer (P2P'09)*, Seattle, WA, Sep. 2009, pp. 99–100.


















# Federated Learning algorithm

---

## Federated Learning (for the Server)

---

Number of iterations:  $T$

Number of clients:  $K$

Fraction of clients per iteration:  $C$

```

1 for each  $t = 1 \dots T$  do
2    $m \leftarrow \max(C \cdot K, 1)$ 
3    $S_t \leftarrow \text{selectRandomSubset}(m, \text{clients})$ 
4   for each  $k \in S_t$ 
5      $\text{send}(k, w_t)$ 
6      $w_{t+1}^k \leftarrow \text{receive}(k)$ 
7    $m_t \leftarrow \sum_{k \in S_t} n_k$ 
8    $w_{t+1} \leftarrow \sum_{k \in S_t} \frac{n_k}{m_t} w_{t+1}^k$ 

```

---



---

## Federated Learning (for a Client)

---



Local dataset:  $D_k$

Number of local epochs:  $E$

Learning rate:  $\eta$

```

1  $\text{receive}(w_t)$ 
2 for each  $\text{epoch} = 1 \dots E$  do
3   for  $\text{batch } b \subset D_k$  do
4      $w_t \leftarrow w_t - \eta \cdot \nabla L(w_t; b)$ 
5  $\text{send}(w_t^k)$ 


```

---

---

```
1 loop
2   if node_is_hub do
3     while time.now() < time_start + delta_time do
4       | models_from_nodes.append(receiveModelFromNode())
5       aggregated_model ← average(models_from_nodes)
6       sendAll(hubs_list, aggregated_model)
7       global_model ← average(models_from_other_hubs)
8       send(nodes_list, global_model)
9   else
10    hubs ← chooseRandom(hubs_list, number_hubs_to_send_model)
11    model ← trainModel(model, data)
12    sendAll(hubs, model)
13    for hub in hubs do
14      | global_models.append(receiveGlobalModelFromHub())
15    model ← average(global_models)
```

---





---

PROOFS Algorithm (active thread)

---

initial peer list: **cache**


cache size: **c**

shuffle length: **l**

node address: **p**

```
1 loop
2   wait( $\Delta$ )
3   subset  $\leftarrow$  selectRandomSubset(cache, l)
4   q  $\leftarrow$  selectRandom(subset)
5   subset.remove(q)
6   subset.add(p)
7   send(q, subset)
8   subsetq  $\leftarrow$  receive(q)
9   subsetq.remove(p)
10  subsetq.removeAll(cache)
11  cache  $\leftarrow$  subsetq
```

---



Before and after a shuffling operation. Node 1 sends addresses {itself, 2, 3} to node 4. Node 4 sends back {5, 6, 8}.

[6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.

- The objective is to obtain an overlay network with  $h$  defined hubs, with  $h$  a parameter of the algorithm, and each hub is connected to all the nodes in the networks. The application running on top will be able to take advantage of this overlay to speed up message transmission in the network.
- **Preferential Attachment:** Drawing from the concept pioneered by Barabási and Albert\*, preferential attachment dictates that new connections in the network are established preferentially with nodes possessing a higher number of existing connections.
- **Random Attachment:** Inspired by gossip-based peer sampling algorithms (\*\*), random attachment ensures that nodes maintain connections with a representative and diverse subset of the network.

[4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.

---

\*

[5] A.-L. Barabási, H. Jeong, Z. Néda, E. Ravasz, A. Schubert, and T. Vicsek, “Evolution of the social network of scientific collaborations,” *Physica A: Statistical mechanics and its applications*, vol. 311, no. 3–4, pp. 590–614, 2002.

\*

[6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.

\*

[14] M. Jelasity, S. Voulgaris, R. Guerraoui, A.-M. Kermarrec, and M. Van Steen, “Gossip-based peer sampling,” *ACM Transactions on Computer Systems (TOCS)*, vol. 25, no. 3, p. 8–es, 2007.

# Elevator's Algorithm

## Elevator Algorithm (active thread)

```

initial peer list: cache
cache size: c
desired number of hubs: h
initial backward list: backward_peers (empty)
1 loop
2   wait( $\Delta$ )
3   frequency_map  $\leftarrow \{\}$ 
4   for peer in cache do
5     peer_cache  $\leftarrow$  send(CACHE_REQUEST, peer)
6     frequency_map  $\leftarrow$  frequency_map  $\cup$  peer_cache
7   preferred  $\leftarrow$  frequency_map.sort().select(h)
8   preferred_backward  $\leftarrow \{\}$ 
9   for peer in preferred do
10    peer_backward_peer  $\leftarrow$  send(BACKWARD_REQUEST,
11    peer)
12    preferred_backward.add(peer_backward_peer)
13   cache  $\leftarrow \{\}$ 
14   cache  $\leftarrow$  preferred  $\cup$  preferred_backward
15   backward_peers  $\leftarrow \{\}$ 


```

## Elevator Algorithm (background thread)

```

1 loop
2   request, peer  $\leftarrow$  receive()
3   if request == CACHE_REQUEST then
4     send(cache, peer)
5     backward_peers.add(peer)
6   if request == BACKWARD_REQUEST
7     then
8       send(randomValue(backward_peers),
9       peer)

```



[15] J. Huang, L. Kong, G. Chen, Q. Xiang, X. Chen, and X. Liu, "Blockchain-Based Federated Learning: A Systematic Survey," *IEEE Network*, vol. 37, no. 6, pp. 150–157, 2023.

**Peer sampling**


**Peer discovery**

**Data request**

**Membership  
management**

**Topology  
management**

**Information  
dissemination**



**Definition.** (Partial View) The **partial view** of  $v$ , denoted  $P(v)$ , is the set of nodes  $v$  can send messages to,  $P(v) \subseteq V$ , with  $|P(v)| \leq c$ ,  $c \ll N$ , and  $P(v) = \text{neigh}_G(v)$ .