

Peer-to-peer protocols for efficient and resilient decentralized learning


Mohamed Amine Legheraba

Supervisors: Maria Potop-Butucaru, Sébastien Tixeul

LIP6 · Sorbonne Université

September 7, 2026


Context



Supervised Classification

Training

Data




Dog



Cat



(data, label)



Model

Labels


Prediction



Trained Model

Cat

Federated Learning



Server



```
1 loop
2   received_models  $\leftarrow$  collect(clients)
3   global_model  $\leftarrow$  merge(received_models)
4   broadcast(global_model)
```

Client


```
1 loop
2   local_model  $\leftarrow$  train(model, local_data)
3   send(local_model, server)
4   model  $\leftarrow$  receive(global_model)
```

[1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Artificial intelligence and statistics*, 2017, pp. 1273–1282.

Structured vs Unstructured networks



Gossip Learning



Gossip Learning (main thread)


ML model: **model**

List of neighbors : **cache**

```
1 loop
2   | Wait for  $\Delta$  time units
3   |  $\text{peer} \leftarrow \text{selectRandom}(\text{cache})$ 
4   |  $\text{send}(\text{peer}, \text{model})$ 
```

[2] R. Ormándi, I. Hegedűs, and M. Jelasity, “Gossip learning with linear models on fully distributed data,” *Concurrency and Computation: Practice and Experience*, vol. 25, no. 4, pp. 556–571, 2013.

Gossip Learning vs Federated Learning




[3] I. Hegedűs, G. Danner, and M. Jelasity, “Decentralized learning works: An empirical comparison of gossip learning and federated learning,” *Journal of Parallel and Distributed Computing*, vol. 148, pp. 109–124, 2021.

State of the Art Summary

	Decentralized	Local data	Fast convergence	Fault-tolerant	No overlay overhead
Centralized learning	✗	✗	✓	✗	✓
Federated learning	✗	✓	✓	✗	✓
Decentralized learning (structured net.)	✓	✓	✓	✗	✗
Gossip learning	✓	✓	✗	✓	✓
Local learning	✓	✓	✗ no generalizing model	✓	✓

No single solution is fully satisfactory



Outline

- **Elevator** — overlay (hub election)
- **Lift** — malicious resilience
- **HEAL** — model aggregation
- **FLAIR** — wireless networks adaptation

Elevator


Definition. (Hub) Let $G = (V, E)$ be the directed overlay graph, where each node $v \in V$ keeps a partial view $C(v) \subseteq V$.

A node $h \in V$ is a **hub** if it appears in the partial view of every node:


$$\forall v \in V, \quad h \in C(v)$$

[4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.


1. One node is more connected





2. It attracts more connections



3. Snowball effect: it becomes a hub



[5] A.-L. Barabási, H. Jeong, Z. Néda, E. Ravasz, A. Schubert, and T. Vicsek, “Evolution of the social network of scientific collaborations,” *Physica A: Statistical mechanics and its applications*, vol. 311, no. 3–4, pp. 590–614, 2002.



Before and after a shuffling operation. Node 1 sends addresses {itself, 2, 3} to node 4. Node 4 sends back {5,6, 8}.

[6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.

Preferential attachment makes hubs emerge...

Random attachment keeps the network resilient...

Can we combine both to get **hubs** that are
resilient?

Mechanism

1. Each node asks its **neighbours** for **their** neighbour lists.
2. It connects to the **h most frequent** nodes.
3. It asks these h hubs for extra incoming connections, to fill a **random subset** of its view.

Observed

- After a few cycles, **hubs emerge** (the preferred h are the same for everyone);
- Hubs are **elected at random** among the network nodes;
- Even after **hub failures**, new nodes are elected as hubs.



Example of topology generated





Figure 1: Indegree distribution of a network generated with Elevator, with 1000 nodes and 10 hubs

Definition. (Stability) Once Elevator has converged to a set of h hubs, both the list of hubs and their number h remain constant (with high probability) over time.


Formally, after convergence time T :

$$\Pr[\forall t \geq T, H(t) = H(T) \wedge \forall v \in V, C_v(t)[1 : h] = H(t)] = 1$$

Definition. (Convergence) The network converges to a stable state containing **exactly** h hubs shared by all nodes.

Idea of the proof.

- A new hub is produced with **non-zero probability** at each cycle.
- The number of hubs **cannot decrease** (until it reaches h).
- Once at least one hub exists, the network is **strongly connected** (w.h.p.) and produces a new hub.



PeerSim was **forked and substantially rewritten** for this thesis:

- Migrated from **SVN to Git**, build rewritten in **Gradle**;
- **Docker + GitLab CI/CD**;
- **Parallelised** the cycle-based engine;
- Failure models from scratch (**crash, churn, Byzantine**);

Parameters

- Network size: **$N = 1000$ nodes**;
- **1000 cycles**, repeated **100 times**;
- Initial topology: **random k -out graph** with **$k = c = 20$** ;
- **$h = 10$ hubs**;
- e.g. brutal crash: **50% of nodes** disconnected at cycle **500**.

Overlay

- In-degree distribution
- Clustering coefficient
- Average path length
- Diameter
- Convergence time

Machine learning



- Accuracy
- Convergence time

Overlay



- **Crash failures**
- **Churn** (dynamic joins / leaves)
- **Byzantine failures**

Machine learning




- **Privacy attacks**
- **Poisoning attacks**

Note: ML-level failures are not addressed in this thesis



- **Go + libp2p** over TCP/IP streams (+ light HTTP per node);
- networks up to **100 nodes**;
- modes: **synchronous / externally-synchronized / asynchronous**;
- **hubs emerge within the first few cycles**
 - matching theory & simulation.



Lift


Coordinated attack — lying about neighbors

addresses of all colluding nodes: **colluders**

backward list: **backward_peers**

```
1 loop
2   (request, peer) ← receive()
3   if request == CACHE_REQUEST
4     backward_peers.add(peer)
5     colluders.shuffle()
6     modified_cache ← colluders[0:c]
7     send(modified_cache, peer)
```


Each colluding node answers with a fake cache that only contains **other colluding nodes**, inflating their visibility.



After convergence, correct nodes **deterministically** re-draw the hubs from a shared random seed — attackers can no longer keep their hub position (effective against infiltration at low collusion rates).


Hub redistribution

```
current hub list: H
network size: N
target hubs: h
1 seed ← sort(H)
2 prng ← Random(seed)
3 selected ← {}
4 while selected.size() < h
5   | id ← prng.nextInt(N)
6   | if id not in selected
7   |   | selected ← selected ∪ {id}
8 H ← selected
```




HEAL





1) Each node train a local model

2) Each node sends its model to a (random) hub



3) Hubs receive and aggregate models


4) Hubs compute a global model together


5) Hubs send back the global model to the nodes


[7] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeuil, “HEAL: Resilient and Self-* Hub-based Learning,” in *International Conference on Advanced Information Networking and Applications*, 2025, pp. 200–211.

Context of experiments

- **Peer-to-Peer:** On the PeerSim simulator.
 - Each generated network had 100 nodes, and the Hub-based topology had 5 hubs
 - In addition to the crash-free context, we ran the algorithm during a context of crash, a context of churn and a context of an attack on the hubs
- **Decentralized Learning:** On the gossip simulator
 - **ML models:** Logistic Regression & LeNet5
 - **Datasets:** Spambase and MNIST
 - Comparison with Federated Learning, Gossip Learning, Epidemic Learning, GAIA, Chord-based Learning, Fedlay








FLAIR






FLAIR operates in **rounds** (LEACH-style clustering):

1. Each node computes its **resource score** (CPU, RAM, GPU, bandwidth).
2. Each round, a fraction of nodes **elect themselves as cluster-heads** (CHs)
— the role rotates.
3. Nodes send their **local models** to their nearest CH, which aggregates them.
4. CHs **change every round** → the aggregated models keep **mixing across clusters**.

- Static network (100 nodes): **highest accuracy** ≈ 0.91 (C-FL / HEAL / Gossip ≈ 0.90 , Gaia ≈ 0.88)
- Convergence up to **2.5x faster** than Gaia
- Resilient to **permanent / temporary / random crashes** (up to 90% nodes)
- Tested under **mobility** and in a **smart farming** scenario



Conclusion



	Decentralized	Fast convergence	Fault-tolerant	Resistant to colluding	Learning	Network type
Elevator	✓	✓	✓	✗	—	P2P
Elevator + Lift	✓	✓	✓	✓	—	P2P
Elevator + HEAL	✓	✓	✓	✗	✓	P2P
FLAIR	✓	✓	✓	✗	✓	Wireless

- **Simple ML models.** Benchmarks are modest — logistic regression on Spambase, LeNet5 on MNIST. No large-scale task, so convergence and accuracy conclusions remain to be confirmed on more ambitious models.
- **Combinatorial explosion.** Network parameters (N, c, h), failure/churn scenarios, fraction of malicious nodes, ML hyper-parameters, and seeds define an enormous space — only a small fraction can be explored.

- **Capability-aware hub election** (heterogeneous nodes) and sensitivity to the initial topology.
- Optimise Elevator (message complexity, memory, simplicity).
- Address **data heterogeneity** in HEAL (personalisation, locally-weighted aggregation).
- A unified **Python simulator** for networking + ML.

This work has already begun during a **3-month internship at NII (Tokyo)**.

- Better robustness at the Overlay layer
- Byzantine robustness in the learning layer (poisoning & privacy attacks)
- Extend HEAL to unsupervised and reinforcement learning
- Vertical federated learning (different features per participant).
- Formal convergence guarantees for HEAL.

- [8] M. A. Legheraba, M. Potop-Butucaru, and S. Tixeul, “Brief Announcement: A Self-* and Persistent Hub Sampling Service,” in *International Symposium on Stabilizing, Safety, and Security of Distributed Systems*, 2024, pp. 461–465.
- [4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeul, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.
- [9] M. A. Legheraba, N. Rachdi, M. Potop-Butucaru, and S. Tixeul, “LIFT: Byzantine Resilient Hub-Sampling,” in *2025 Thirteenth International Symposium on Computing and Networking (CANDAR)*, 2025, pp. 1–9. — **Outstanding Paper Award**
- [7] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeul, “HEAL: Resilient and Self-* Hub-based Learning,” in *International Conference on Advanced Information Networking and Applications*, 2025, pp. 200–211.
- [10] I. Boutebicha, B. Zaghdoudi, M. A. Legheraba, and M. P. Butucaru, “FLAIR: Distributed Federated Learning with Dynamic Clustering,” in *Proceedings of the 14th International Conference on Networked Systems (NETYS 2026)*, in Lecture Notes in Computer Science. Agadir, Morocco: Springer, 2026.
- [11] M. A. Legheraba, M. Potop-Butucaru, S. Tixeul, and S. Fdida, “Des noeuds anonymes hissés au rang de stars,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- [12] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeul, “Les étoiles montantes de l'apprentissage distribué,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- A journal extension (journal version of this work) has been **submitted to the IEEE/ACM Transactions on Networking**; we are awaiting the **final review** (currently in **minor revision**).

Bibliography

- [1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Artificial intelligence and statistics*, 2017, pp. 1273–1282.
- [2] R. Ormándi, I. Hegedűs, and M. Jelasity, “Gossip learning with linear models on fully distributed data,” *Concurrency and Computation: Practice and Experience*, vol. 25, no. 4, pp. 556–571, 2013.
- [3] I. Hegedűs, G. Danner, and M. Jelasity, “Decentralized learning works: An empirical comparison of gossip learning and federated learning,” *Journal of Parallel and Distributed Computing*, vol. 148, pp. 109–124, 2021.
- [4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.
- [5] A.-L. Barabási, H. Jeong, Z. Néda, E. Ravasz, A. Schubert, and T. Vicsek, “Evolution of the social network of scientific collaborations,” *Physica A: Statistical mechanics and its applications*, vol. 311, no. 3–4, pp. 590–614, 2002.
- [6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.
- [7] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeuil, “HEAL: Resilient and Self-* Hub-based Learning,” in *International Conference on Advanced Information Networking and Applications*, 2025, pp. 200–211.

- [8] M. A. Legheraba, M. Potop-Butucaru, and S. Tixeuil, “Brief Announcement: A Self-* and Persistent Hub Sampling Service,” in *International Symposium on Stabilizing, Safety, and Security of Distributed Systems*, 2024, pp. 461–465.
- [9] M. A. Legheraba, N. Rachdi, M. Potop-Butucaru, and S. Tixeuil, “LIFT: Byzantine Resilient Hub-Sampling,” in *2025 Thirteenth International Symposium on Computing and Networking (CANDAR)*, 2025, pp. 1–9.
- [10] I. Boutebicha, B. Zaghdoudi, M. A. Legheraba, and M. P. Butucaru, “FLAIR: Distributed Federated Learning with Dynamic Clustering,” in *Proceedings of the 14th International Conference on Networked Systems (NETYS 2026)*, in *Lecture Notes in Computer Science*. Agadir, Morocco: Springer, 2026.
- [11] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Des noeuds anonymes hissés au rang de stars,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- [12] M. A. Legheraba, S. Galkiewicz, M. Potop-Butucaru, and S. Tixeuil, “Les étoiles montantes de l'apprentissage distribué,” in *ALGOTEL 2025–27èmes Rencontres Francophones sur les Aspects Algorithmiques des Télécommunications*, 2025.
- [13] A. Montresor and M. Jelasity, “PeerSim: A Scalable P2P Simulator,” in *Proc. of the 9th Int. Conference on Peer-to-Peer (P2P'09)*, Seattle, WA, Sep. 2009, pp. 99–100.
- [14] M. Jelasity, S. Voulgaris, R. Guerraoui, A.-M. Kermarrec, and M. Van Steen, “Gossip-based peer sampling,” *ACM Transactions on Computer Systems (TOCS)*, vol. 25, no. 3, p. 8-es, 2007.


- [15] J. Huang, L. Kong, G. Chen, Q. Xiang, X. Chen, and X. Liu, “Blockchain-Based Federated Learning: A Systematic Survey,” *IEEE Network*, vol. 37, no. 6, pp. 150–157, 2023.



Appendix

- **Elevator**: a self-organising overlay that makes hubs emerge — provides a **structured overlay** for structured aggregation without any coordinator.
- **HEAL**: federated learning directly onto the structured overlay — recovers the efficiency of FL while staying decentralised.
- **Lift**: hardens hub election against colluding byzantines (up to 10%).
- **FLAIR**: validates the modular architecture in wireless networks (LEACH-style clustering).



Answer: structured aggregation is possible **without centralisation**, if the overlay layer is designed with that goal in mind.



Federated Learning (FedAvg)

```
1 for  $t = 0$  to  $T - 1$  do
2   server broadcasts  $\theta^{(t)}$ 
   for each client  $i$  in parallel:
3    $\theta_i \leftarrow$  local training (lr  $\eta$ ,  $E$ 
   steps)
   server aggregates:  $\theta^{(t+1)} \leftarrow$ 
4    $\sum w_i \theta_i$ 
5 end for
```

[1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Artificial intelligence and statistics*, 2017, pp. 1273–1282.



HEAL Learning: The Hub Algorithm

```
duration to wait for model: delta_time
list of all hubs: hubs_list
1 nb_hubs  $\leftarrow$  hubs_list.size()
2 loop
3   models  $\leftarrow$  {}
4   time  $\leftarrow$  time.now()
5   backwards_list  $\leftarrow$  {}
6   while time.now() < time + delta_time do
7     (peer_model, peer)  $\leftarrow$  receive()
8     models.append(peer_model)
9     backwards_list.append(peer)
10  average_model  $\leftarrow$  average(models)
11  send(hubs_list, average_model)
12  models_hubs  $\leftarrow$  {}
13  models_hubs.append(average_model)
14  nb_receive  $\leftarrow$  0
15  while nb_receive < nb_hubs - 1 do
16    hub_model  $\leftarrow$  receive()
17    nb_receive  $\leftarrow$  nb_receive + 1
18    models_hubs.append(hub_model)
19  global_model  $\leftarrow$  average(models_hubs)
20  send(backwards_list, global_model)
```

HEAL Learning: The Client Algorithm

```
duration to wait for model: delta_time
The ML model, initialized at random: model
The local data: data
list of all hubs: hubs_list
number of hubs to send the model: s
1 loop
2   model  $\leftarrow$  trainModel(model, data)
3   hubs  $\leftarrow$  chooseRandom(hubs_list, s)
4   for hub in hubs do
5     | send(hub, model)
6   hubs_models  $\leftarrow$  {}
7   for hub in hubs do
8     | model_hub  $\leftarrow$  receive()
9     | hubs_models.append(model_hub)
10  model  $\leftarrow$  average(hubs_models)
```

Elevator Algorithm (background thread)

```
1 loop
2   request, peer  $\leftarrow$  receive()
3   if request == CACHE_REQUEST then
4     send(cache, peer)
5     backward_peers.add(peer)
6   if request == BACKWARD_REQUEST then
7     send(backward_peers, peer)
```



- Simulations done with the Java PeerSim^{*} simulator (modified),
with the cycle based mode
- Comparaisons against 3 peer sampling algorithms : Newscast, Proofs and Phenix (Power-law)



	Value
Network size	1000
Number of cycles for each simulation	1000
Number of times each simulation was run (with different seed)	100
c parameter (cache size)	20
h parameter (number of hubs)	10



- All simulations were run on 16 vCPU, using 64G of memory.
- Code is available at <https://gitlab.lip6.fr/legheraba/elevator>



*



[13] A. Montresor and M. Jelasity, "PeerSim: A Scalable P2P Simulator," in *Proc. of the 9th Int. Conference on Peer-to-Peer (P2P'09)*, Seattle, WA, Sep. 2009, pp. 99–100.














Federated Learning algorithm

Federated Learning (for the Server)

Number of iterations: T

Number of clients: K

Fraction of clients per iteration: C

```

1 for each  $t = 1 \dots T$  do
2    $m \leftarrow \max(C \cdot K, 1)$ 
3    $S_t \leftarrow \text{selectRandomSubset}(m, \text{clients})$ 
4   for each  $k \in S_t$ 
5      $\text{send}(k, w_t)$ 
6      $w_{t+1}^k \leftarrow \text{receive}(k)$ 
7    $m_t \leftarrow \sum_{k \in S_t} n_k$ 
8    $w_{t+1} \leftarrow \sum_{k \in S_t} \frac{n_k}{m_t} w_{t+1}^k$ 

```

Federated Learning (for a Client)



Local dataset: D_k

Number of local epochs: E


Learning rate: η

```

1  $\text{receive}(w_t)$ 
2 for each  $\text{epoch} = 1 \dots E$  do
3   for  $\text{batch } b \subset D_k$  do
4      $w_t \leftarrow w_t - \eta \cdot \nabla L(w_t; b)$ 
5  $\text{send}(w_t^k)$ 

```

```
1 loop
2   if node_is_hub do
3     while time.now() < time_start + delta_time do
4       | models_from_nodes.append(receiveModelFromNode())
5       aggregated_model ← average(models_from_nodes)
6       sendAll(hubs_list, aggregated_model)
7       global_model ← average(models_from_other_hubs)
8       send(nodes_list, global_model)
9   else
10    hubs ← chooseRandom(hubs_list, number_hubs_to_send_model)
11    model ← trainModel(model, data)
12    sendAll(hubs, model)
13    for hub in hubs do
14      | global_models.append(receiveGlobalModelFromHub())
15    model ← average(global_models)
```



PROOFS Algorithm (active thread)


initial peer list: **cache**

cache size: **c**

shuffle length: **l**

node address: **p**

```
1 loop
2   wait( $\Delta$ )
3   subset  $\leftarrow$  selectRandomSubset(cache, l)
4   q  $\leftarrow$  selectRandom(subset)
5   subset.remove(q)
6   subset.add(p)
7   send(q, subset)
8   subsetq  $\leftarrow$  receive(q)
9   subsetq.remove(p)
10  subsetq.removeAll(cache)
11  cache  $\leftarrow$  subsetq
```



Before and after a shuffling operation. Node 1 sends addresses {itself, 2, 3} to node 4. Node 4 sends back {5, 6, 8}.

[6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.

- The objective is to obtain an overlay network with h defined hubs, with h a parameter of the algorithm, and each hub is connected to all the nodes in the networks. The application running on top will be able to take advantage of this overlay to speed up message transmission in the network.
- **Preferential Attachment:** Drawing from the concept pioneered by Barabási and Albert*, preferential attachment dictates that new connections in the network are established preferentially with nodes possessing a higher number of existing connections.
- **Random Attachment:** Inspired by gossip-based peer sampling algorithms (**), random attachment ensures that nodes maintain connections with a representative and diverse subset of the network.

[4] M. A. Legheraba, M. Potop-Butucaru, S. Tixeuil, and S. Fdida, “Emergent Peer-to-Peer Multi-Hub Topology,” in *2024 22nd International Symposium on Network Computing and Applications (NCA)*, 2024, pp. 219–226.

*

[5] A.-L. Barabási, H. Jeong, Z. Néda, E. Ravasz, A. Schubert, and T. Vicsek, “Evolution of the social network of scientific collaborations,” *Physica A: Statistical mechanics and its applications*, vol. 311, no. 3–4, pp. 590–614, 2002.

*

[6] A. Stavrou, D. Rubenstein, and S. Sahu, “A lightweight, robust p2p system to handle flash crowds,” in *10th IEEE International Conference on Network Protocols, 2002. Proceedings.*, 2002, pp. 226–235.

*

[14] M. Jelasity, S. Voulgaris, R. Guerraoui, A.-M. Kermarrec, and M. Van Steen, “Gossip-based peer sampling,” *ACM Transactions on Computer Systems (TOCS)*, vol. 25, no. 3, p. 8–es, 2007.

Elevator's Algorithm

Elevator Algorithm (active thread)

```

initial peer list: cache
cache size: c
desired number of hubs: h
initial backward list: backward_peers (empty)
1 loop
2   wait( $\Delta$ )
3   frequency_map  $\leftarrow \{\}$ 
4   for peer in cache do
5     | peer_cache  $\leftarrow$  send(CACHE_REQUEST, peer)
6     | frequency_map  $\leftarrow$  frequency_map  $\cup$  peer_cache
7   preferred  $\leftarrow$  frequency_map.sort().select(h)
8   preferred_backward  $\leftarrow \{\}$ 
9   for peer in preferred do
10    | peer_backward_peer  $\leftarrow$  send(BACKWARD_REQUEST,
11    | peer)
12    | preferred_backward.add(peer_backward_peer)
13   cache  $\leftarrow \{\}$ 
14   cache  $\leftarrow$  preferred  $\cup$  preferred_backward
14   backward_peers  $\leftarrow \{\}$ 


```

Elevator Algorithm (background thread)

```

1 loop
2   request, peer  $\leftarrow$  receive()
3   if request == CACHE_REQUEST then
4     | send(cache, peer)
5     | backward_peers.add(peer)
6   if request == BACKWARD_REQUEST
7   then
8     | send(randomValue(backward_peers),
9     | peer)

```



[15] J. Huang, L. Kong, G. Chen, Q. Xiang, X. Chen, and X. Liu, "Blockchain-Based Federated Learning: A Systematic Survey," *IEEE Network*, vol. 37, no. 6, pp. 150–157, 2023.

Peer sampling


Peer discovery

Data request

**Membership
management**

**Topology
management**

**Information
dissemination**




Definition. (Partial View) The **partial view** of v , denoted $P(v)$, is the set of nodes v can send messages to, $P(v) \subseteq V$, with $|P(v)| \leq c$, $c \ll N$, and $P(v) = \text{neigh}_G(v)$.

Spambase + Logistic regression

Variables Table		
Variable Name	Role	Type
word_freq_make	Feature	Continuous
word_freq_address	Feature	Continuous
word_freq_all	Feature	Continuous
word_freq_3d	Feature	Continuous
word_freq_our	Feature	Continuous
word_freq_over	Feature	Continuous
word_freq_remove	Feature	Continuous

4601 emails, binary
57 parameters

MNIST + LeNet5



70k digits, 10 classes
CNN, 60k parameters